

Lecture 17: The Conjugate Gradient Method

COMP5930M Scientific Computation

Today

Outline

Gradient descent

Conjugate directions

Conjugate gradient method

Examples

Next

Reference

*An Introduction to the Conjugate Gradient Method,
Without the Agonizing Pain*

Jonathan Shewchuk

<http://www.cs.cmu.edu/~jrs/jrspapers.html#cg>

- ▶ A well-written, intuitive description
- ▶ This lecture (mostly) follows this paper

Outline

We can develop the algorithm as a series of enhancements, starting from a form of the Gradient Descent algorithm.

- ▶ The Gradient-Descent Algorithm
- ▶ The Conjugate Directions Algorithm
- ▶ The Conjugate Gradient Method

1. The problem

We want to compute the n -dimensional vector \mathbf{x} satisfying $\mathbf{Ax} = \mathbf{b}$, where \mathbf{A} is a large, sparse matrix

We will assume \mathbf{A} is symmetric and *positive-definite* (SPD)

- ▶ SPD means that for any vector $\mathbf{p} \neq \mathbf{0}$, scalar $s = \mathbf{p}^T \mathbf{A} \mathbf{p} > 0$

Recall we can also define a set of n eigenvalues λ_i and eigenvectors \mathbf{e}_i for symmetric \mathbf{A}

- ▶ For SPD matrices $\lambda_i > 0$ for all i

The quadratic form

Consider a related problem:

- Find \mathbf{x} that minimises the scalar function $f(\mathbf{x})$

$$f(\mathbf{x}) = \frac{1}{2} \mathbf{x}^T \mathbf{A} \mathbf{x} - \mathbf{b}^T \mathbf{x} + c$$

with some scalar constant c

- At a minimum of f we have the necessary condition

$$\frac{\partial f}{\partial \mathbf{x}} = \mathbf{A} \mathbf{x} - \mathbf{b} = 0$$

hence the minimum point \mathbf{x} also solves $\mathbf{A} \mathbf{x} = \mathbf{b}$

Minimising $f(\mathbf{x})$

- ▶ Conjugate Gradient (CG) and related algorithms were designed as **minimisation algorithms**
- ▶ SPD matrices guarantee the function $f(\mathbf{x})$ is *strictly convex* and hence that the point where the gradient is zero is a minimum

Eigenvalues/vectors

- ▶ They are the primary *analysis* tool for these algorithms but we do not need to compute them (general eigendecomposition would require $\mathcal{O}(N^3)$)
- ▶ The shape of the space we search is related to the eigenvalues and eigenvectors: $\mathbf{A}\mathbf{u}_i = \lambda_i\mathbf{u}_i$
- ▶ For each pair λ_i, \mathbf{u}_i the search space is stretched
 - ▶ by a factor $1/\lambda_i$
 - ▶ in the direction \mathbf{u}_i

History

- ▶ CG was designed in the 1960s as a *direct* solution algorithm
 - ▶ Formally it will terminate in n steps at the exact solution
 - ▶ It was discarded as inefficient compared to standard direct algorithms
- ▶ In the 1970s it was reinvented as an iterative algorithm
 - ▶ Due to the search process it will often be **close** to a solution in less than n steps

The residual

- Recall we defined the residual \mathbf{r}_i for a given approximate solution \mathbf{x}_i

$$\mathbf{r}_i = \mathbf{b} - \mathbf{A}\mathbf{x}_i$$

- \mathbf{r}_i defines a local search direction of [steepest-descent](#). Why?

$$f\mathbf{x}(\mathbf{x}_i) = \mathbf{A}\mathbf{x}_i - \mathbf{b} = -\mathbf{r}_i$$

2. The Gradient-Descent Algorithm

Define an update

$$\mathbf{x}_{i+1} = \mathbf{x}_i + \alpha_i \mathbf{r}_i$$

for some scalar step length α_i

- ▶ We should compute the optimal distance α_i for that search direction
- ▶ This is the **line-search** problem, (considered in Lecture 6, but now for a linear system)
- ▶ We can **exactly** derive (see Ref p6)

$$\alpha_i = \frac{\mathbf{r}_i^T \mathbf{r}_i}{\mathbf{r}_i^T \mathbf{A} \mathbf{r}_i}$$

Orthogonality

- ▶ If we minimise the gradient in one direction, \mathbf{r}_i , the next step of steepest-descent must be orthogonal
- ▶ This is equivalent to observing that

$$\mathbf{r}_i^T \mathbf{r}_{i+1} = 0$$

- ▶ A series of steps defines a series of searches in consecutively orthogonal directions
- ▶ We cannot guarantee that all directions \mathbf{r}_i will be orthogonal to all other search directions $\mathbf{r}_j \perp \mathbf{r}_i$ for all $i \neq j$.

3. Conjugate Directions Algorithm

Define an update

$$\mathbf{x}_{i+1} = \mathbf{x}_i + \alpha_i \mathbf{d}_i$$

for some scalar step length α_i

- ▶ An improved search algorithm that guarantees the set of search directions, \mathbf{d}_i , are **mutually orthogonal**, for any $i \neq j$

$$\mathbf{d}_i^T \mathbf{d}_j = 0$$

- ▶ For example, in Cartesian space try $\mathbf{d}_i = \{\mathbf{e}_x, \mathbf{e}_y, \mathbf{e}_z, \dots\}$ (called **cyclic coordinate search**)
 - ▶ Any point in space can be located by successively searching each coordinate direction

In practice?

To have orthogonal search directions \mathbf{d}_i we must require that at each step the error has to be orthogonal to the search direction, $\mathbf{d}_i^T \mathbf{e}_{i+1}$.

- We can derive the optimal step size α_i as

$$\alpha_i = -\frac{\mathbf{d}_i^T \mathbf{e}_i}{\mathbf{d}_i^T \mathbf{d}_i}$$

for any orthogonal set \mathbf{d}_i

- This requires the error $\mathbf{e}_i = \mathbf{x} - \mathbf{x}_i$ and is **not computable**

A-orthogonality

- ▶ Note: when \mathbf{A} is symmetric positive-definite, we can define the \mathbf{A} -norm of a vector using the quadratic form

$$\|\mathbf{x}\|_A := \sqrt{\mathbf{x}^T \mathbf{A} \mathbf{x}}.$$

- ▶ \mathbf{A} -orthogonality of directions \mathbf{d}_i is defined by

$$\mathbf{d}_i^T \mathbf{A} \mathbf{d}_j = 0$$

for any $i \neq j$

- ▶ For this set we can derive a computable form

$$\alpha_i = \frac{\mathbf{d}_i^T \mathbf{A} \mathbf{e}_i}{\mathbf{d}_i^T \mathbf{A} \mathbf{d}_i} = \frac{\mathbf{d}_i^T \mathbf{r}_i}{\mathbf{d}_i^T \mathbf{A} \mathbf{d}_i}$$

- ▶ Note, that choosing $\mathbf{d}_i = \mathbf{r}_i$ would produce a modified gradient-descent algorithm

Computing \mathbf{d}_i

At step i , we use the residual \mathbf{r}_i as a search direction, but enforce \mathbf{A} -orthogonality by subtracting the previous search directions :

$$\mathbf{d}_i = \mathbf{r}_i + \sum_{k=0}^{i-1} \beta_{ik} \mathbf{d}_k = \mathbf{r}_i - \sum_{k=0}^{i-1} \text{proj}_{\mathbf{d}_k}(\mathbf{r}_i)$$

where $\text{proj}_{\mathbf{d}_k}(\mathbf{r}_i)$ is the projection of \mathbf{r}_i to the direction \mathbf{d}_k given by

$$\beta_{ik} = -\frac{\mathbf{r}_i^T \mathbf{A} \mathbf{d}_k}{\mathbf{d}_k^T \mathbf{A} \mathbf{d}_k}$$

- Although appearing complex we are simply subtracting out components not \mathbf{A} -orthogonal to a previous search direction

Note

- ▶ This algorithm is generally called **Gram-Schmidt**
 - ▶ In general it has $\mathcal{O}(n^2)$ expense, as the work increases with i
- ▶ The set of search directions at each step i is called a **Krylov subspace** for A

4. The Conjugate Gradient Method

- ▶ The choice of basing the search direction \mathbf{d}_i on the residual \mathbf{r}_i allows a significant amount of algebraic simplification
- ▶ We can show
 - ▶ The residual \mathbf{r}_i is \mathbf{A} -orthogonal to every previous search direction \mathbf{d}_j
 - ▶ $\beta_{ik} = 0$ unless $i = k + 1$
- ▶ The expense of CG is then $\mathcal{O}(nz)$

The CG algorithm

- ▶ $\mathbf{d}_0 = \mathbf{r}_0 = \mathbf{b} - \mathbf{A}\mathbf{x}_0$
- ▶ $i = 0, 1, 2, \dots$
 - ▶ $\alpha = \frac{\mathbf{r}_i^T \mathbf{r}_i}{\mathbf{d}_i^T \mathbf{A} \mathbf{d}_i}$
 - ▶ $\mathbf{x}_{i+1} = \mathbf{x}_i + \alpha \mathbf{d}_i$
 - ▶ $\mathbf{r}_{i+1} = \mathbf{r}_i - \alpha \mathbf{A} \mathbf{d}_i$
 - ▶ $\beta = \frac{\mathbf{r}_{i+1}^T \mathbf{r}_{i+1}}{\mathbf{r}_i^T \mathbf{r}_i}$
 - ▶ $\mathbf{d}_{i+1} = \mathbf{r}_{i+1} + \beta \mathbf{d}_i$

Note: Only one matrix-multiply required per iteration

Efficiency

- ▶ For efficiency we require the approximate solution in a minimum number of iterations
 - ▶ we can relate the convergence to the distribution of eigenvalues of the matrix
 - ▶ the specific relationship depends on the iterative scheme
- ▶ The reference paper is more comprehensive and outlines the mathematical derivations
- ▶ The CG algorithm is often described more formally without *physical* insight

Examples

- ▶ Matlab code `runPCG.m` available on VLE
- ▶ Uses Matlab library `pcg.m` function
- ▶ Matrix generated with the Matlab `gallery()` function

Improving the Conjugate Gradient algorithm

- ▶ General convergence result for CG:

$$\|\mathbf{e}_k\|_A \leq 2 \left(\frac{\sqrt{\kappa(\mathbf{A})} - 1}{\sqrt{\kappa(\mathbf{A})} + 1} \right)^k \|\mathbf{e}_0\|_A.$$

- ▶ We can improve the convergence,
ie. reduce the number of iterations required,
through **preconditioning**
 - ▶ We solve $\mathbf{M}^{-1}\mathbf{Ax} = \mathbf{M}^{-1}\mathbf{b}$
 - ▶ The matrix $\mathbf{M}^{-1}\mathbf{A}$ has an improved eigenvalue distribution
 - ▶ This implies a reduced Condition Number κ

Summary

- ▶ Conjugate gradient (CG) method efficient iterative solver for **symmetric, positive-definite** matrices
- ▶ Method based on construction of search directions \mathbf{d}_i that are **A-orthogonal** (belong in Krylov subspace), followed by descent method with exact line-search
- ▶ For non-symmetric problems we have other Krylov-methods:
 - ▶ Generalised minimal residual -method (GMRES)
 - ▶ Biconjugate gradient stabilised -method (BiCGSTAB)
- ▶ Number of iterations depends on condition number $\kappa(\mathbf{A})$